

```

#import the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

data = pd.read_csv("dataset.csv")
print("Dataset Shape:", data.shape)

Dataset Shape: (250, 8)

print(data.columns)

Index(['Age', 'Annual Income', 'Spending Score', 'Gender', 'Location',
       'Family Size', 'Occupation', 'Education'],
      dtype='object')

print(data.size)

2000

#data preprocessing using label encoding
from sklearn.preprocessing import LabelEncoder, StandardScaler

obj=LabelEncoder()
data['Gender']=obj.fit_transform(data['Gender'])
data['Location']=obj.fit_transform(data['Location'])
data['Occupation']=obj.fit_transform(data['Occupation'])
data['Education']=obj.fit_transform(data['Education'])

# Data preprocessing using label encoding
# le = LabelEncoder()
# categorical_columns = ['Gender', 'Location', 'Occupation',
# 'Education']
# for col in categorical_columns:
#     data[col] = le.fit_transform(data[col])

# Feature scaling
scaler = StandardScaler()
features_scaled = scaler.fit_transform(data)
features_df = pd.DataFrame(features_scaled, columns=data.columns)

features_df

```

	Age	Annual Income	Spending Score	Gender	Location	Family
Size \						
0	1.368290	0.359492	-1.239574	-1.0	-1.074789	-
0.254442						
1	0.577005	-0.960610	1.496569	1.0	-1.074789	-
0.254442						
2	-0.530794	0.780383	-1.543590	-1.0	-1.074789	
1.480387						
3	-1.084694	1.514264	-0.530204	1.0	0.930415	-

```

1.410994
4    -0.056023    -1.472113    1.226333    -1.0    0.930415
0.323835
...      ...      ...      ...      ...      ...
...
245    1.210033    -0.382408    -1.138236    1.0    0.930415
1.480387
246    0.814391    -1.637155    1.124994    -1.0    -1.074789
1.480387
247   -1.005565    -0.653573    1.260112    -1.0    0.930415    -
0.832718
248    1.526547    -0.646659    0.618301    -1.0    0.930415
0.323835
249    1.051776    -1.151583    -0.563983    1.0    -1.074789    -
0.254442

```

```

      Occupation  Education
0      -0.090085  -0.686878
1     -1.113772   1.415811
2      0.933603   0.714914
3      0.933603   0.714914
4      1.445447  -1.387775
...      ...      ...
245    -1.113772  -1.387775
246    -0.090085  -1.387775
247    -1.113772   0.014018
248    -0.090085   1.415811
249    -0.090085  -1.387775

```

```
[250 rows x 8 columns]
```

```
features_df.columns
```

```
Index(['Age', 'Annual Income', 'Spending Score', 'Gender', 'Location',
      'Family Size', 'Occupation', 'Education'],
      dtype='object')
```

```
features_df.size
```

```
2000
```

```
features_df.shape
```

```
(250, 8)
```

```
features_df.head(10)
```

```
features_df.tail(4)
```

```
features_df.sample(7)
```

```

      Age  Annual Income  Spending Score  Gender  Location  Family
Size \


```

30	-0.530794	0.427776	-1.138236	-1.0	0.930415	-
	0.254442					
55	-1.242951	-0.121038	0.348064	1.0	-1.074789	-
	1.410994					
171	0.181363	-0.200654	-0.935558	-1.0	-1.074789	-
	1.410994					
34	-0.056023	-0.851221	-1.577370	1.0	-1.074789	
	1.480387					
4	-0.056023	-1.472113	1.226333	-1.0	0.930415	
	0.323835					
95	1.605676	-1.044548	1.226333	-1.0	-1.074789	-
	0.254442					
207	-0.451665	1.457601	-1.138236	1.0	0.930415	-
	0.832718					

	Occupation	Education
30	-0.090085	1.415811
55	-1.113772	-0.686878
171	-1.625616	-0.686878
34	-0.601928	1.415811
4	1.445447	-1.387775
95	0.421759	-0.686878
207	-1.625616	-0.686878

*#feature selection*

```
import seaborn as sns
plt.figure(figsize=(12,6))
sns.heatmap(data.corr(),annot=True)
plt.title("using correlation matrix ")
Text(0.5, 1.0, 'using correlation matrix ')
```



```
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering, Birch, SpectralClustering

selector = SelectKBest(score_func=f_classif, k=4)
X_selected = selector.fit_transform(features_df, data['Spending Score'])
selected_features_mask = selector.get_support()
selected_feature_names = features_df.columns[selected_features_mask]
print("Most significant features:", selected_feature_names)


Most significant features: Index(['Annual Income', 'Spending Score', 'Location', 'Family Size'], dtype='object')

c:\Users\lsrin\anaconda3\Lib\site-packages\sklearn\feature_selection\_univariate_selection.py:113: RuntimeWarning: divide by zero encountered in divide
  f = msb / msw

# 2ndPCA
pca = PCA()
X_pca = pca.fit_transform(features_df)
explained_variance = pca.explained_variance_ratio_
cumulative_variance = np.cumsum(explained_variance)
print(len(explained_variance))
```

8

```
plt.figure(figsize=(10, 6))
r=range(1, len(explained_variance) + 1) #1,9 inclusive,exclusive
plt.plot(r, cumulative_variance)
plt.xlabel('no of Components')
plt.ylabel('Cumulative EVR')
plt.title('PCA-CumulativeEVR')
Text(0.5, 1.0, 'PCA-CumulativeEVR')
```



```
from sklearn.feature_selection import SelectKBest, f_classif
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans, AgglomerativeClustering, Birch,
SpectralClustering

selector = SelectKBest(score_func=f_classif, k=4)
X_selected = selector.fit_transform(features_df, data['Spending
Score'])
selected_features_mask = selector.get_support()
selected_feature_names = features_df.columns[selected_features_mask]
print("Most significant features:", selected_feature_names)
```

```
Most significant features: Index(['Annual Income', 'Spending Score', 'Location', 'Family Size'], dtype='object')
```

```
c:\Users\lsrin\anaconda3\Lib\site-packages\sklearn\feature_selection\_univariate_selection.py:113: RuntimeWarning: divide by zero encountered in divide
```

```
    f = msb / msw
```

```
# 2ndPCA
```

```
pca = PCA()
```

```
X_pca = pca.fit_transform(features_df)
```

```
explained_variance = pca.explained_variance_ratio_ #variances blw the pc explained "how"
```

```
cumulative_variance = np.cumsum(explained_variance) #total variances
```

```
print(len(explained_variance))
```

```
8
```

```
plt.figure(figsize=(10, 6))
```

```
r=range(1, len(explained_variance) + 1) #1,9 inclusive,exclusive
```

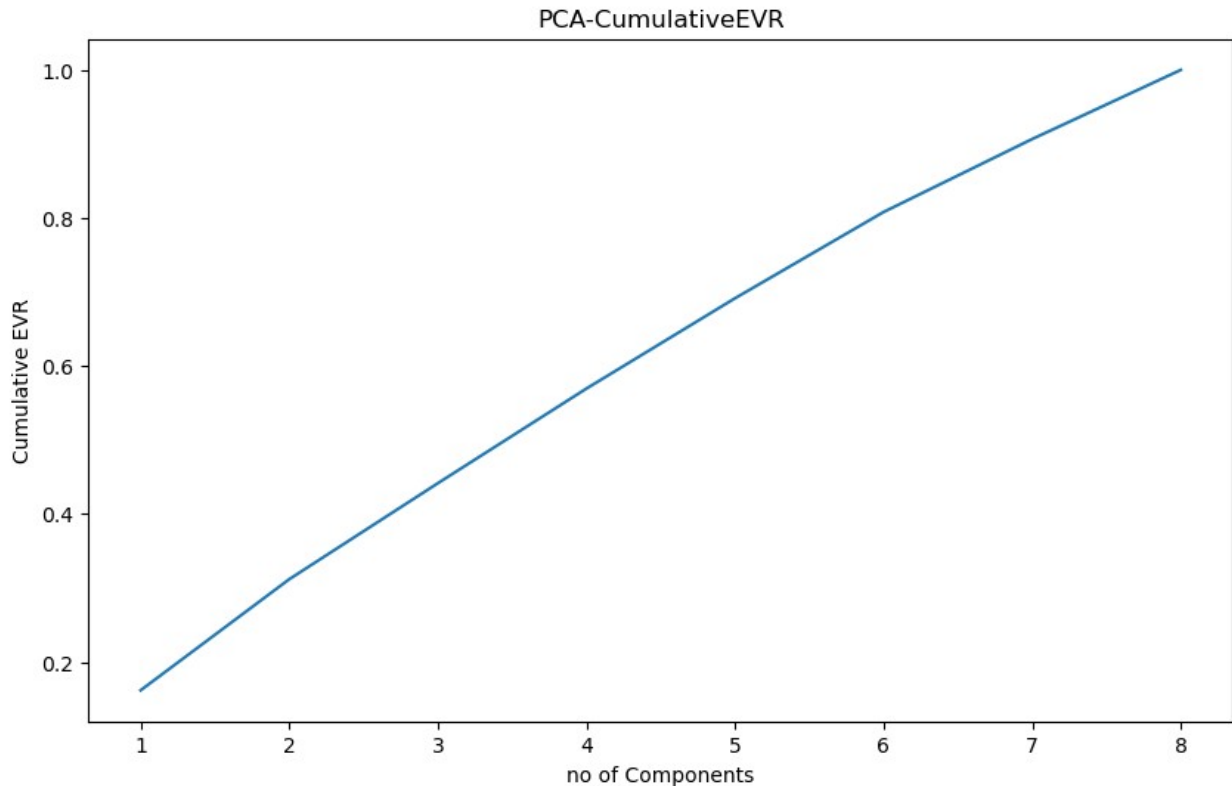
```
plt.plot(r, cumulative_variance)
```

```
plt.xlabel('no of Components')
```

```
plt.ylabel('Cumulative EVR')
```

```
plt.title('PCA-CumulativeEVR')
```

```
Text(0.5, 1.0, 'PCA-CumulativeEVR')
```



```
# no of cluster analysis
from sklearn.cluster import KMeans, AgglomerativeClustering, Birch
from sklearn.metrics import silhouette_score
import scipy.cluster.hierarchy as sch

# K-Means for different cluster numbers
K = range(2, 11) # testing clusters from 2 to 10
inertias = [] # for K-Means elbow methodfor within clustering we can use ss
silhouette_scores = [] # for measuring cluster quality

for k in K:
    kmeans = KMeans(n_clusters=k, random_state=42)
    kmeans_labels = kmeans.fit_predict(features_df)
    inertias.append(kmeans.inertia_)
    silhouette_scores.append(silhouette_score(features_df,
kmeans_labels))

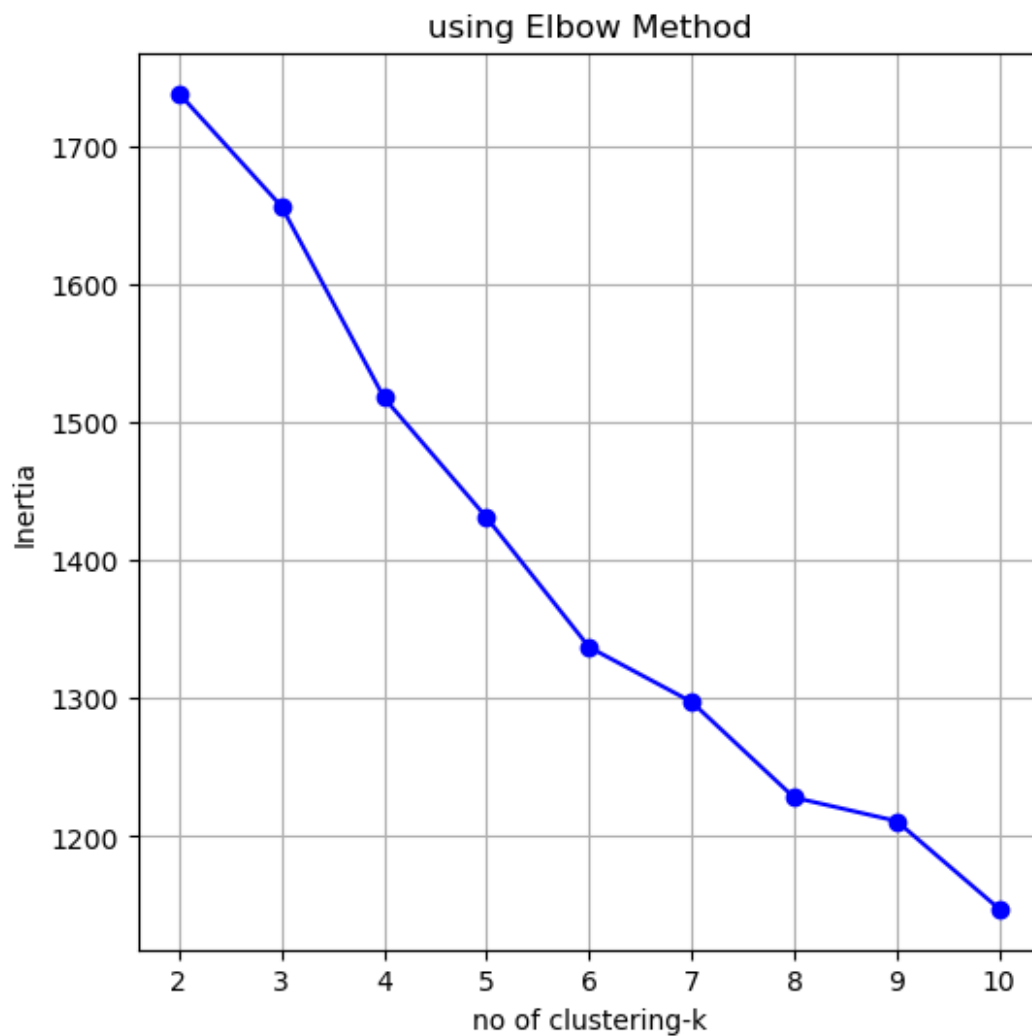
c:\Users\lsrin\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1446: UserWarning: KMeans is known to have a memory leak on
Windows with MKL, when there are less chunks than available threads.
You can avoid it by setting the environment variable
OMP_NUM_THREADS=1.
  warnings.warn(
c:\Users\lsrin\anaconda3\Lib\site-packages\sklearn\cluster\
_kmeans.py:1446: UserWarning: KMeans is known to have a memory leak on
```






```
plt.xlabel('no of clustering-k')
plt.ylabel('Inertia')
plt.title('using Elbow Method')
plt.grid(True)
plt.show()

#Silhouette Analysis
plt.figure(figsize=(5,5))
plt.plot(K, silhouette_scores, 'ro-')
plt.xlabel('no of clustering (k)')
plt.ylabel('Silhouette-score')
plt.title('Silhouette-Analysis')
plt.grid(True)
plt.show()
```





```
#apply clustering with optimal k  
k_optimal = 3
```

```
#apply different clustering methods  
kmeans = KMeans(n_clusters=k_optimal, random_state=42)  
hierarchical = AgglomerativeClustering(n_clusters=k_optimal)  
birch = Birch(n_clusters=k_optimal)
```

```
print(kmeans)  
print(hierarchical)  
print(birch)
```

```
KMeans(n_clusters=3, random_state=42)  
AgglomerativeClustering(n_clusters=3)  
Birch()
```

```
#cluster labels  
kmeans_labels = kmeans.fit_predict(features_df)  
hierarchical_labels = hierarchical.fit_predict(features_df)  
birch_labels = birch.fit_predict(features_df)
```

```
c:\Users\lsrin\anaconda3\Lib\site-packages\sklearn\cluster\  
_kmeans.py:1446: UserWarning: KMeans is known to have a memory leak on
```

Windows with MKL, when there are less chunks than available threads.  
You can avoid it by setting the environment variable  
OMP\_NUM\_THREADS=1.  
warnings.warn(

*# evaluate and compare methods*

```
def evaluate_clustering(data, labels, method_name):  
    score = silhouette_score(data, labels)  
    print(f"{method_name} Silhouette Score: {score:.3f}")  
  
print("clustering quality evaluation:")  
evaluate_clustering(features_df, kmeans_labels, "K-Means")  
evaluate_clustering(features_df, hierarchical_labels, "Hierarchical")  
evaluate_clustering(features_df, birch_labels, "BIRCH")
```

```
clustering quality evaluation:  
K-Means Silhouette Score: 0.096  
Hierarchical Silhouette Score: 0.095  
BIRCH Silhouette Score: 0.083
```

*#visual*

```
pca = PCA(n_components=2)  
X_pca = pca.fit_transform(features_df)
```

*# K-Means*

```
plt.figure(figsize=(10, 6))  
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=kmeans_labels, cmap='viridis')  
plt.title('K-Means Clustering')  
plt.xlabel('First Principal Component')  
plt.ylabel('Second Principal Component')  
plt.colorbar(label='Cluster Labels')  
plt.show()
```

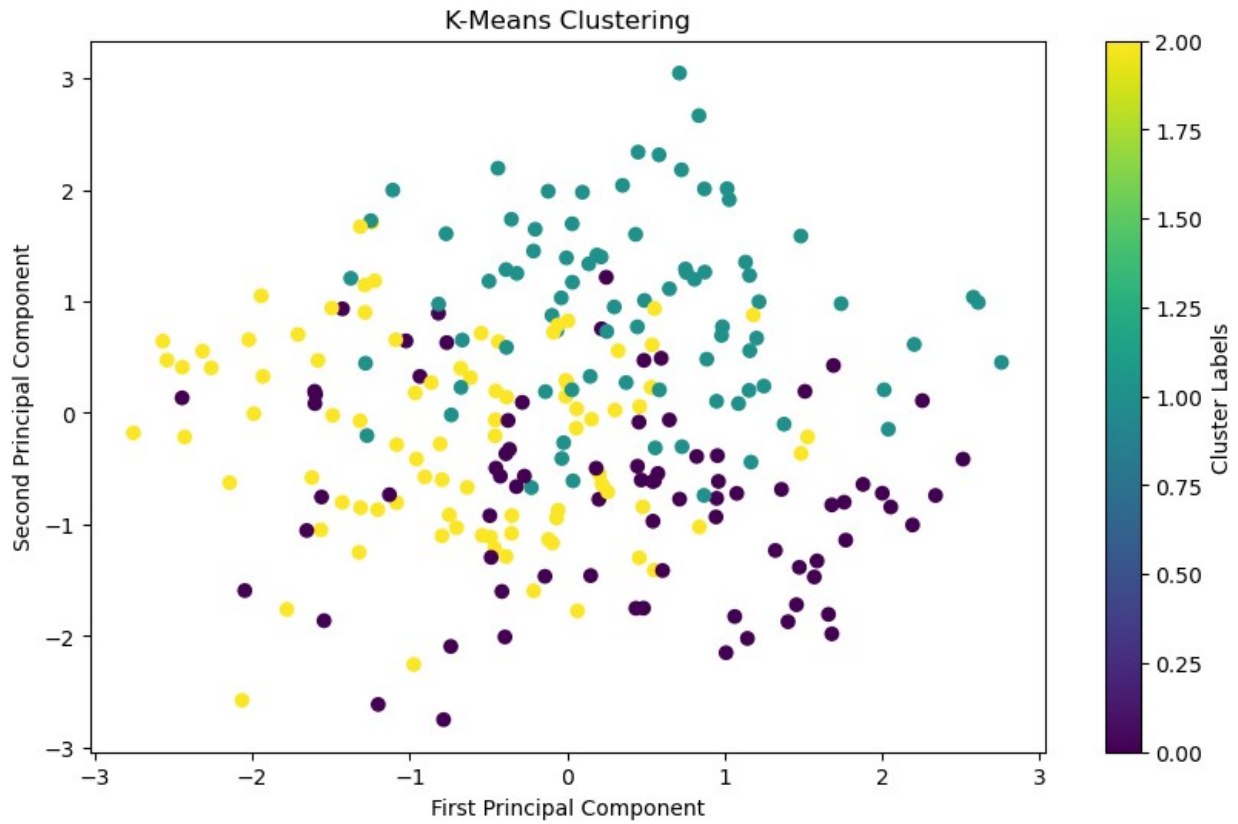
*#hierarchical*

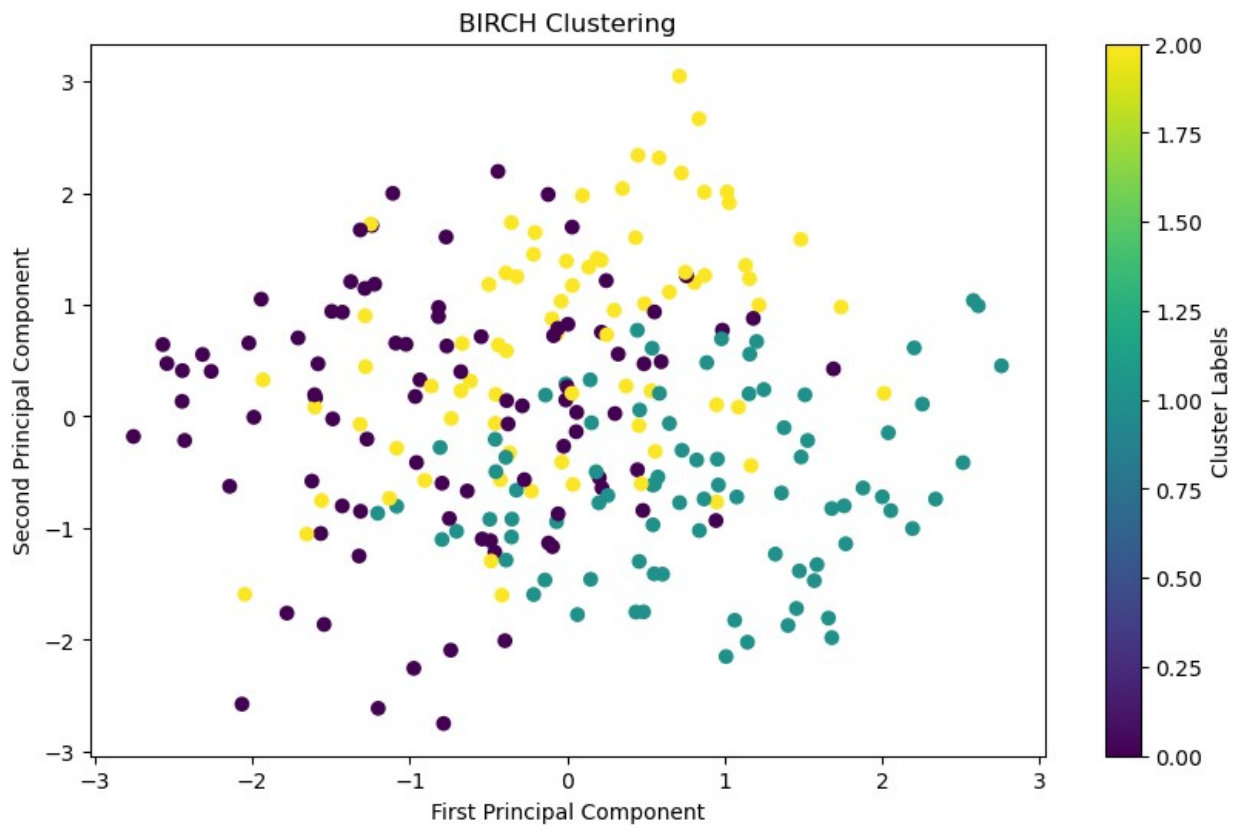
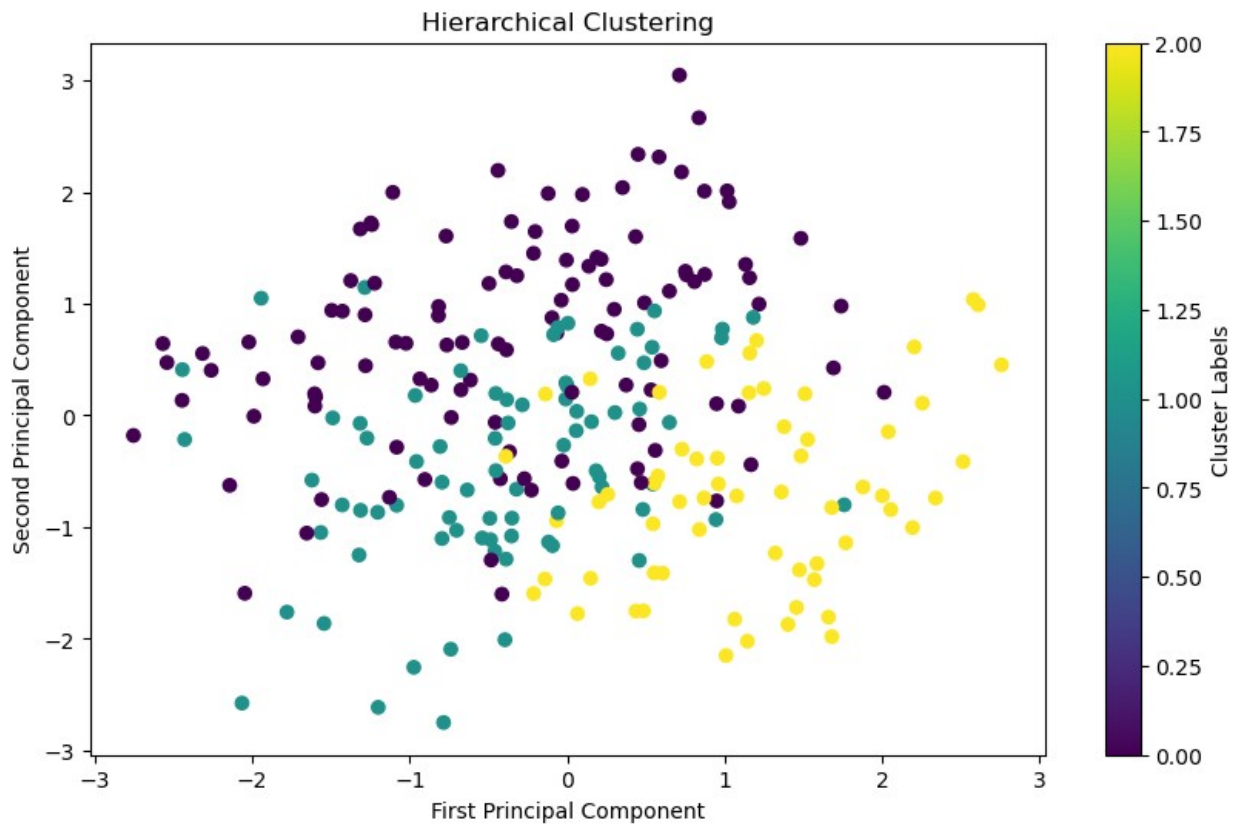
```
plt.figure(figsize=(10, 6))  
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=hierarchical_labels,  
cmap='viridis')  
plt.title('Hierarchical Clustering')  
plt.xlabel('First Principal Component')  
plt.ylabel('Second Principal Component')  
plt.colorbar(label='Cluster Labels')  
plt.show()
```

*# BIRCH*

```
plt.figure(figsize=(10, 6))  
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=birch_labels, cmap='viridis')  
plt.title('BIRCH Clustering')
```

```
plt.xlabel('First Principal Component')  
plt.ylabel('Second Principal Component')  
plt.colorbar(label='Cluster Labels')  
plt.show()
```





```

#evaluation Metrics
from sklearn.metrics import calinski_harabasz_score,
davies_bouldin_score

def evaluate_clustering_detailed(data, labels, method_name):
    silhouette = silhouette_score(data, labels)
    calinski = calinski_harabasz_score(data, labels)
    davies = davies_bouldin_score(data, labels)

    print(f"\nDetailed Evaluation for {method_name}:")
    print(f"Silhouette Score: {silhouette:.3f}")
    print(f"Calinski-Harabasz Score: {calinski:.3f}")
    print(f"Davies-Bouldin Score: {davies:.3f}")
    return silhouette, calinski, davies

# evaluate each method with detailed metrics
results = {}
for method, labels in [("K-Means", kmeans_labels),
                        ("Hierarchical", hierarchical_labels),
                        ("BIRCH", birch_labels)]:
    results[method] = evaluate_clustering_detailed(features_df,
labels, method)

# Print comparative analysis
print("\nComparative Analysis:")
print("Based on the evaluation metrics:")
print("1.silhouette Score: Higher is better (range: -1 to 1)")
print("2.calinski-Harabasz Score: Higher is better")
print("3.davies-Bouldin Score: Lower is better")

# Determine best method
best_silhouette = max(results.items(), key=lambda x: x[1][0])
print("presenting the list best_silhouette :",best_silhouette)
print(f"\nbest performing method based on Silhouette score:
{best_silhouette[0]}")

##conclusion part
print("1.feature Selection identified the following key features:",
list(selected_feature_names))
print("2.optimal number of clusters based on Elbow method:",
k_optimal)
print("3.best performing clustering method:", best_silhouette[0])

```

```

Detailed Evaluation for K-Means:
Silhouette Score: 0.096
Calinski-Harabasz Score: 25.607
Davies-Bouldin Score: 2.557

```

Detailed Evaluation for Hierarchical:

Silhouette Score: 0.095

Calinski-Harabasz Score: 26.549

Davies-Bouldin Score: 2.536

Detailed Evaluation for BIRCH:

Silhouette Score: 0.083

Calinski-Harabasz Score: 23.497

Davies-Bouldin Score: 2.749

Comparative Analysis:

Based on the evaluation metrics:

1.silhouette Score: Higher is better (range: -1 to 1)

2.calinski-Harabasz Score: Higher is better

3.davies-Bouldin Score: Lower is better

presenting the list best\_silhouette : ('K-Means',  
(0.09586237556192426, 25.607180550441758, 2.5568902549578847))

best performing method based on Silhouette score: K-Means

1.feature Selection identified the following key features: ['Annual  
Income', 'Spending Score', 'Location', 'Family Size']

2.optimal number of clusters based on Elbow method: 3

3.best performing clustering method: K-Means