

```
import numpy as np;
import pandas as pd
import matplotlib.pyplot as plt;
import seaborn as sn

#load the dataset
dataset=pd.read_csv("c:\\Users\\lsrin\\Downloads\\TS-2\\Adavance_ML\\
climate.csv")
print(dataset)
```

	STATION	DATE	REPORT_TYPE	SOURCE
BackupElements \				
0	72518014735	2015-01-01T23:59:00	SOD	6
PRECIP				
1	72518014735	2015-01-02T23:59:00	SOD	6
PRECIP				
2	72518014735	2015-01-03T23:59:00	SOD	6
PRECIP				
3	72518014735	2015-01-04T23:59:00	SOD	6
PRECIP				
4	72518014735	2015-01-05T23:59:00	SOD	6
PRECIP				
...
...				
2663	72518014735	2022-05-27T23:59:00	SOD	6
PRECIP				
2664	72518014735	2022-05-28T23:59:00	SOD	6
PRECIP				
2665	72518014735	2022-05-29T23:59:00	SOD	6
PRECIP				
2666	72518014735	2022-05-30T23:59:00	SOD	6
PRECIP				
2667	72518014735	2022-05-31T23:59:00	SOD	6
PRECIP				

	BackupElevation	BackupEquipment	BackupLatitude	BackupLongitude
\				
0	260	PLASTIC	42.6918	-73.83109
1	260	PLASTIC	42.6918	-73.83109
2	260	PLASTIC	42.6918	-73.83109
3	260	PLASTIC	42.6918	-73.83109
4	260	PLASTIC	42.6918	-73.83109
...
2663	260	PLASTIC	42.6812	-73.81650

2664	260	PLASTIC	42.6812	-73.81650
2665	260	PLASTIC	42.6812	-73.81650
2666	260	PLASTIC	42.6812	-73.81650
2667	260	PLASTIC	42.6812	-73.81650
BackupName ... DailyPeakWindDirection DailyPeakWindSpeed				
\				
0	NWS ALBANY, NY	...	190.0	26.0
1	NWS ALBANY, NY	...	250.0	30.0
2	NWS ALBANY, NY	...	170.0	21.0
3	NWS ALBANY, NY	...	290.0	33.0
4	NWS ALBANY, NY	...	280.0	42.0
...
2663	NWS ALBANY, NY	...	160.0	28.0
2664	NWS ALBANY, NY	...	310.0	26.0
2665	NWS ALBANY, NY	...	90.0	13.0
2666	NWS ALBANY, NY	...	200.0	15.0
2667	NWS ALBANY, NY	...	250.0	29.0
DailyPrecipitation DailySnowDepth DailySnowfall \				
0	0.00	0.0	0.0	
1	T	0.0	T	
2	0.57	0.0	1.6	
3	0.22	1.0	0.0	
4	T	0.0	T	
...	
2663	0.00	0	0.0	
2664	0.04	0	0.0	
2665	0.00	0	0.0	
2666	0.00	0	0.0	
2667	0.00	0	0.0	
DailySustainedWindDirection DailySustainedWindSpeed Sunrise				
Sunset \				
0		190.0	20.0	726.0

1632.0			
1	310.0	23.0	726.0
1633.0			
2	160.0	15.0	726.0
1634.0			
3	290.0	24.0	726.0
1635.0			
4	290.0	32.0	726.0
1636.0			
...
...			
2663	160.0	21.0	423.0
1922.0			
2664	310.0	22.0	422.0
1923.0			
2665	180.0	9.0	421.0
1924.0			
2666	190.0	12.0	421.0
1925.0			
2667	250.0	21.0	420.0
1926.0			

	WindEquipmentChangeDate
0	2006-09-08
1	2006-09-08
2	2006-09-08
3	2006-09-08
4	2006-09-08
...	...
2663	2006-09-08
2664	2006-09-08
2665	2006-09-08
2666	2006-09-08
2667	2006-09-08

[2668 rows x 32 columns]

dataset.columns

```
Index(['STATION', 'DATE', 'REPORT_TYPE', 'SOURCE', 'BackupElements',
      'BackupElevation', 'BackupEquipment', 'BackupLatitude',
      'BackupLongitude', 'BackupName',
      'DailyAverageDewPointTemperature',
      'DailyAverageDryBulbTemperature',
      'DailyAverageRelativeHumidity',
      'DailyAverageSeaLevelPressure', 'DailyAverageStationPressure',
      'DailyAverageWetBulbTemperature', 'DailyAverageWindSpeed',
      'DailyCoolingDegreeDays',
      'DailyDepartureFromNormalAverageTemperature',
      'DailyHeatingDegreeDays', 'DailyMaximumDryBulbTemperature',
```

```

'DailyMinimumDryBulbTemperature', 'DailyPeakWindDirection',
'DailyPeakWindSpeed', 'DailyPrecipitation', 'DailySnowDepth',
'DailySnowfall', 'DailySustainedWindDirection',
'DailySustainedWindSpeed', 'Sunrise', 'Sunset',
'WindEquipmentChangeDate'],
dtype='object')

```

#DATA PROCESSING

```

dataset.head(10)
dataset.tail(8)
dataset.sample(7)
print(dataset.columns)
print(dataset.dtypes)
print(dataset.shape[0]) #column
print(dataset.shape[1]) #rows
print(dataset.size)

```

```

Index(['STATION', 'DATE', 'REPORT_TYPE', 'SOURCE', 'BackupElements',
      'BackupElevation', 'BackupEquipment', 'BackupLatitude',
      'BackupLongitude', 'BackupName',
      'DailyAverageDewPointTemperature',
      'DailyAverageDryBulbTemperature',
      'DailyAverageRelativeHumidity',
      'DailyAverageSeaLevelPressure', 'DailyAverageStationPressure',
      'DailyAverageWetBulbTemperature', 'DailyAverageWindSpeed',
      'DailyCoolingDegreeDays',
      'DailyDepartureFromNormalAverageTemperature',
      'DailyHeatingDegreeDays', 'DailyMaximumDryBulbTemperature',
      'DailyMinimumDryBulbTemperature', 'DailyPeakWindDirection',
      'DailyPeakWindSpeed', 'DailyPrecipitation', 'DailySnowDepth',
      'DailySnowfall', 'DailySustainedWindDirection',
      'DailySustainedWindSpeed', 'Sunrise', 'Sunset',
      'WindEquipmentChangeDate'],
      dtype='object')

```

STATION	int64
DATE	object
REPORT_TYPE	object
SOURCE	int64
BackupElements	object
BackupElevation	int64
BackupEquipment	object
BackupLatitude	float64
BackupLongitude	float64
BackupName	object
DailyAverageDewPointTemperature	float64
DailyAverageDryBulbTemperature	float64
DailyAverageRelativeHumidity	float64
DailyAverageSeaLevelPressure	float64
DailyAverageStationPressure	float64
DailyAverageWetBulbTemperature	float64

DailyAverageWindSpeed	float64
DailyCoolingDegreeDays	float64
DailyDepartureFromNormalAverageTemperature	float64
DailyHeatingDegreeDays	float64
DailyMaximumDryBulbTemperature	float64
DailyMinimumDryBulbTemperature	float64
DailyPeakWindDirection	float64
DailyPeakWindSpeed	float64
DailyPrecipitation	object
DailySnowDepth	object
DailySnowfall	object
DailySustainedWindDirection	float64
DailySustainedWindSpeed	float64
Sunrise	float64
Sunset	float64
WindEquipmentChangeDate	object
dtype: object	
2668	
32	
85376	

dataset.isnull().sum()

STATION	0
DATE	0
REPORT_TYPE	0
SOURCE	0
BackupElements	0
BackupElevation	0
BackupEquipment	0
BackupLatitude	0
BackupLongitude	0
BackupName	0
DailyAverageDewPointTemperature	0
DailyAverageDryBulbTemperature	0
DailyAverageRelativeHumidity	0
DailyAverageSeaLevelPressure	0
DailyAverageStationPressure	0
DailyAverageWetBulbTemperature	0
DailyAverageWindSpeed	0
DailyCoolingDegreeDays	0
DailyDepartureFromNormalAverageTemperature	0
DailyHeatingDegreeDays	0
DailyMaximumDryBulbTemperature	0
DailyMinimumDryBulbTemperature	0
DailyPeakWindDirection	0
DailyPeakWindSpeed	0
DailyPrecipitation	0
DailySnowDepth	0
DailySnowfall	0

```
DailySustainedWindDirection      0
DailySustainedWindSpeed          0
Sunrise                         0
Sunset                          0
WindEquipmentChangeDate         0
dtype: int64
```

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2668 entries, 0 to 2667
```

```
Data columns (total 32 columns):
```

#	Column	Non-Null Count	Dtype
0	STATION	2668 non-null	int64
1	DATE	2668 non-null	object
2	REPORT_TYPE	2668 non-null	object
3	SOURCE	2668 non-null	int64
4	BackupElements	2668 non-null	object
5	BackupElevation	2668 non-null	int64
6	BackupEquipment	2668 non-null	object
7	BackupLatitude	2668 non-null	float64
8	BackupLongitude	2668 non-null	float64
9	BackupName	2668 non-null	object
10	DailyAverageDewPointTemperature	2668 non-null	float64
11	DailyAverageDryBulbTemperature	2668 non-null	float64
12	DailyAverageRelativeHumidity	2668 non-null	float64
13	DailyAverageSeaLevelPressure	2668 non-null	float64
14	DailyAverageStationPressure	2668 non-null	float64
15	DailyAverageWetBulbTemperature	2668 non-null	float64
16	DailyAverageWindSpeed	2668 non-null	float64

17	DailyCoolingDegreeDays	2668	non-null
float64			
18	DailyDepartureFromNormalAverageTemperature	2668	non-null
float64			
19	DailyHeatingDegreeDays	2668	non-null
float64			
20	DailyMaximumDryBulbTemperature	2668	non-null
float64			
21	DailyMinimumDryBulbTemperature	2668	non-null
float64			
22	DailyPeakWindDirection	2668	non-null
float64			
23	DailyPeakWindSpeed	2668	non-null
float64			
24	DailyPrecipitation	2668	non-null
object			
25	DailySnowDepth	2668	non-null
object			
26	DailySnowfall	2668	non-null
object			
27	DailySustainedWindDirection	2668	non-null
float64			
28	DailySustainedWindSpeed	2668	non-null
float64			
29	Sunrise	2668	non-null
float64			
30	Sunset	2668	non-null
float64			
31	WindEquipmentChangeDate	2668	non-null
object			

dtypes: float64(20), int64(3), object(9)
memory usage: 667.1+ KB

```
# Ensure 'DailyPrecipitation' is numeric and fill NaN values with 0
dataset['DailyPrecipitation'] =
pd.to_numeric(dataset['DailyPrecipitation'], errors='coerce')
dataset['DailyPrecipitation'].isnull().sum()
dataset['DailyPrecipitation'].fillna(0)
```

0	0.00
1	0.00
2	0.57
3	0.22
4	0.00
	...
2663	0.00
2664	0.04
2665	0.00
2666	0.00

```
2667    0.00
```

```
Name: DailyPrecipitation, Length: 2668, dtype: float64
```

```
dataset['DailyPrecipitation'].isnull().sum()
```

```
0
```

```
#do some statistical for int64,float64
```

```
dataNumerical=dataset.select_dtypes(include=['int64','float64'])
```

```
#print(dataNumerical.columns)
```

```
print(dataNumerical.dtypes)
```

STATION	int64
SOURCE	int64
BackupElevation	int64
BackupLatitude	float64
BackupLongitude	float64
DailyAverageDewPointTemperature	float64
DailyAverageDryBulbTemperature	float64
DailyAverageRelativeHumidity	float64
DailyAverageSeaLevelPressure	float64
DailyAverageStationPressure	float64
DailyAverageWetBulbTemperature	float64
DailyAverageWindSpeed	float64
DailyCoolingDegreeDays	float64
DailyDepartureFromNormalAverageTemperature	float64
DailyHeatingDegreeDays	float64
DailyMaximumDryBulbTemperature	float64
DailyMinimumDryBulbTemperature	float64
DailyPeakWindDirection	float64
DailyPeakWindSpeed	float64
DailyPrecipitation	float64
DailySustainedWindDirection	float64
DailySustainedWindSpeed	float64
Sunrise	float64
Sunset	float64

```
dtype: object
```

```
#drop the unnecessary attribute
```

```
dataNumerical = dataNumerical.drop(columns=[  
    'STATION', 'SOURCE', 'BackupElevation', 'BackupLatitude',  
    'BackupLongitude', 'Sunrise', 'Sunset'  
])
```

```
dataNumerical.columns
```

```
Index(['DailyAverageDewPointTemperature',  
      'DailyAverageDryBulbTemperature',  
      'DailyAverageRelativeHumidity', 'DailyAverageSeaLevelPressure',  
      'DailyAverageStationPressure',  
      'DailyAverageWetBulbTemperature',
```



```

        'DailyAverageWindSpeed', 'DailyCoolingDegreeDays',
        'DailyDepartureFromNormalAverageTemperature',
'DailyHeatingDegreeDays',
        'DailyMaximumDryBulbTemperature',
'DailyMinimumDryBulbTemperature',
        'DailyPeakWindDirection', 'DailyPeakWindSpeed',
'DailyPrecipitation',
        'DailySustainedWindDirection', 'DailySustainedWindSpeed'],
dtype='object')

```

#Using linear Regression

Features matrix (X) with the temperature column

```
X = dataNumerical[['DailyAverageDryBulbTemperature']]
```

Target vector (y) with precipitation values

```
y = dataNumerical['DailyPrecipitation']
```

Print X and y to verify

```
print(X)
```

```
print(y.isnull().sum())
```

```

      DailyAverageDryBulbTemperature
0                                25.0
1                                32.0
2                                27.0
3                                39.0
4                                27.0
...                               ...
2663                           70.0
2664                           68.0
2665                           66.0
2666                           72.0
2667                           79.0

```

```
[2668 rows x 1 columns]
```

```
0
```

```
from sklearn.model_selection import train_test_split
```

Split the data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=32)
```

Print the number of rows in X_train

```
print(X_train.shape[0])
```

```
2134
```

#train the model

```
from sklearn.linear_model import LinearRegression
```

Initialize and train the linear regression model

```

model = LinearRegression()
model.fit(X_train, y_train)

LinearRegression()

# Make predictions
y_pred = model.predict(X_test)
#print(y_pred)
#y_pred

# Evaluate the model
from sklearn.metrics import mean_squared_error, r2_score
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

Mean Squared Error: 0.0729113990218066
R-squared: 0.0025889789634933047

```

2. Feature Selection Techniques(Low Variance Filter)

The goal is to remove features with very low variance, as they don't provide useful information. For example, if a feature value barely changes across data samples, it won't help the model predict precipitation.

S1: Remove the X feature =DailyPrecipitation

s2:VarianceThreshold with 0.1 which are <0.1 contain low variances s3:fit the selectors

```

from sklearn.feature_selection import VarianceThreshold

# Step 1: Drop the target column (DailyPrecipitation) from the feature matrix
X = dataNumerical.drop(columns=['DailyPrecipitation'])
print("All Features:")
print(X.columns)

# Step 2: Initialize the VarianceThreshold selector with a specific threshold (e.g., 0.1)
selector = VarianceThreshold(threshold=0.1)

# Step 3: Fit the selector on the feature matrix
selector.fit(X)

# Step 4: Calculate variances for each feature
variances = X.var()
print("\nVariance of each feature:")
print(variances)

```

```

# Step 5: Identify threshold value
threshold = selector.threshold
print(f"\nVariance Threshold: {threshold}")

# Step 6: Identify selected and removed features based on the
threshold
selected_columns = variances[variances >= threshold].index
removed_columns = variances[variances < threshold].index

# Step 7: Print selected and removed features
print("\nSelected Features after Low Variance Filter:")
print(selected_columns)

print("\nRemoved Features with Low Variance:")
print(removed_columns)

All Features:
Index(['DailyAverageDewPointTemperature',
'DailyAverageDryBulbTemperature',
'DailyAverageRelativeHumidity', 'DailyAverageSeaLevelPressure',
'DailyAverageStationPressure',
'DailyAverageWetBulbTemperature',
'DailyAverageWindSpeed', 'DailyCoolingDegreeDays',
'DailyDepartureFromNormalAverageTemperature',
'DailyHeatingDegreeDays',
'DailyMaximumDryBulbTemperature',
'DailyMinimumDryBulbTemperature',
'DailyPeakWindDirection', 'DailyPeakWindSpeed',
'DailySustainedWindDirection', 'DailySustainedWindSpeed'],
dtype='object')

Variance of each feature:
DailyAverageDewPointTemperature      365.431000
DailyAverageDryBulbTemperature        351.461615
DailyAverageRelativeHumidity          179.596433
DailyAverageSeaLevelPressure           0.050074
DailyAverageStationPressure            0.048773
DailyAverageWetBulbTemperature        290.359311
DailyAverageWindSpeed                  14.805719
DailyCoolingDegreeDays                 17.913956
DailyDepartureFromNormalAverageTemperature  67.288101
DailyHeatingDegreeDays                 260.312571
DailyMaximumDryBulbTemperature         400.148238
DailyMinimumDryBulbTemperature         328.421192
DailyPeakWindDirection                8249.827997
DailyPeakWindSpeed                     89.043312
DailySustainedWindDirection            8253.098128
DailySustainedWindSpeed                48.192936
dtype: float64

```

Variance Threshold: 0.1

Selected Features after Low Variance Filter:

```
Index(['DailyAverageDewPointTemperature',  
      'DailyAverageDryBulbTemperature',  
        'DailyAverageRelativeHumidity',  
      'DailyAverageWetBulbTemperature',  
        'DailyAverageWindSpeed', 'DailyCoolingDegreeDays',  
        'DailyDepartureFromNormalAverageTemperature',  
      'DailyHeatingDegreeDays',  
        'DailyMaximumDryBulbTemperature',  
      'DailyMinimumDryBulbTemperature',  
        'DailyPeakWindDirection', 'DailyPeakWindSpeed',  
        'DailySustainedWindDirection', 'DailySustainedWindSpeed'],  
      dtype='object')
```

Removed Features with Low Variance:

```
Index(['DailyAverageSeaLevelPressure', 'DailyAverageStationPressure'],  
      dtype='object')
```

```
X_selected = dataNumerical[selected_columns] # This extracts the data matrix
```

```
# Target variable remains the same
```

```
y = dataNumerical['DailyPrecipitation']
```

```
# Verify shapes of X and y
```

```
print(X_selected.shape)
```

```
print(y.shape)
```

```
# Split the data into train and test sets
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X_selected, y,  
                                                    test_size=0.2, random_state=42)
```

```
print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(2668, 14)
```

```
(2668,)
```

```
(2134, 14) (534, 14) (2134,) (534,)
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import r2_score, mean_squared_error
```

```
# Initialize and train the model
```

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

```
# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
r2LVF = r2_score(y_test, y_pred)
mseLVF = mean_squared_error(y_test, y_pred)

print(f'R2 Score: {r2LVF:.2f}')
print(f'Mean Squared Error: {mseLVF:.2f}')

R2 Score: 0.29
Mean Squared Error: 0.05
```

These visualizations will give you:

Feature distributions – Helps understand spread and outliers. Scatter plot – Visualizes relationships between features and target. Correlation heatmap – Identifies multicollinearity. Residual plot – Checks for bias in predictions. Line plot – Visualizes how well predictions align with actual values. Bar plot – Summarizes model metrics.

```
dataNumerical.columns

Index(['DailyAverageDewPointTemperature',
      'DailyAverageDryBulbTemperature',
      'DailyAverageRelativeHumidity', 'DailyAverageSeaLevelPressure',
      'DailyAverageStationPressure',
      'DailyAverageWetBulbTemperature',
      'DailyAverageWindSpeed', 'DailyCoolingDegreeDays',
      'DailyDepartureFromNormalAverageTemperature',
      'DailyHeatingDegreeDays',
      'DailyMaximumDryBulbTemperature',
      'DailyMinimumDryBulbTemperature',
      'DailyPeakWindDirection', 'DailyPeakWindSpeed',
      'DailyPrecipitation',
      'DailySustainedWindDirection', 'DailySustainedWindSpeed'],
      dtype='object')
```

Applying a High Correlation Filter (HCF) is a great way to reduce feature dimensionality by removing features that are highly correlated with each other. This helps to prevent multicollinearity issues in your model, which can affect the interpretability and performance of regression models.

```
X_selected = dataNumerical.drop(columns=['DailyPrecipitation']) #
Ensure to remove the target

# Step 1: Calculate the correlation matrix
correlation_matrix = X_selected.corr()
```

```

# Step 2: Set a correlation threshold
threshold = 0.8

# Step 3: Identify features to remove
correlated_features = set() # Set to store features to remove

# Iterate through the correlation matrix
for i in range(len(correlation_matrix.columns)):
    for j in range(i):
        if abs(correlation_matrix.iloc[i, j]) > threshold:
            # Get the name of the feature
            colname = correlation_matrix.columns[i]
            correlated_features.add(colname)

# Display correlated features
print("Correlated Features to Remove:")
print(correlated_features)

# Step 4: Drop correlated features from the feature matrix
X_filtered = X_selected.drop(columns=correlated_features)

# Display the remaining features
print("\nFinal Features after High Correlation Filter:")
print(X_filtered.columns)

Correlated Features to Remove:
{'DailyMinimumDryBulbTemperature', 'DailyAverageWetBulbTemperature',
'DailyAverageStationPressure', 'DailyHeatingDegreeDays',
'DailySustainedWindSpeed', 'DailyAverageDryBulbTemperature',
'DailyMaximumDryBulbTemperature', 'DailyPeakWindSpeed'}

Final Features after High Correlation Filter:
Index(['DailyAverageDewPointTemperature',
'DailyAverageRelativeHumidity',
'DailyAverageSeaLevelPressure', 'DailyAverageWindSpeed',
'DailyCoolingDegreeDays',
'DailyDepartureFromNormalAverageTemperature',
'DailyPeakWindDirection', 'DailySustainedWindDirection'],
      dtype='object')

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error

# Use the filtered features for training and evaluation
y = dataNumerical['DailyPrecipitation']

# Split the filtered data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_filtered, y,
test_size=0.2, random_state=42)

```

```

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
r2HCF= r2_score(y_test, y_pred)
msehcf= mean_squared_error(y_test, y_pred)


print(f'R2 Score: {r2HCF:.2f}')
print(f'Mean Squared Error: {msehcf:.2f}')

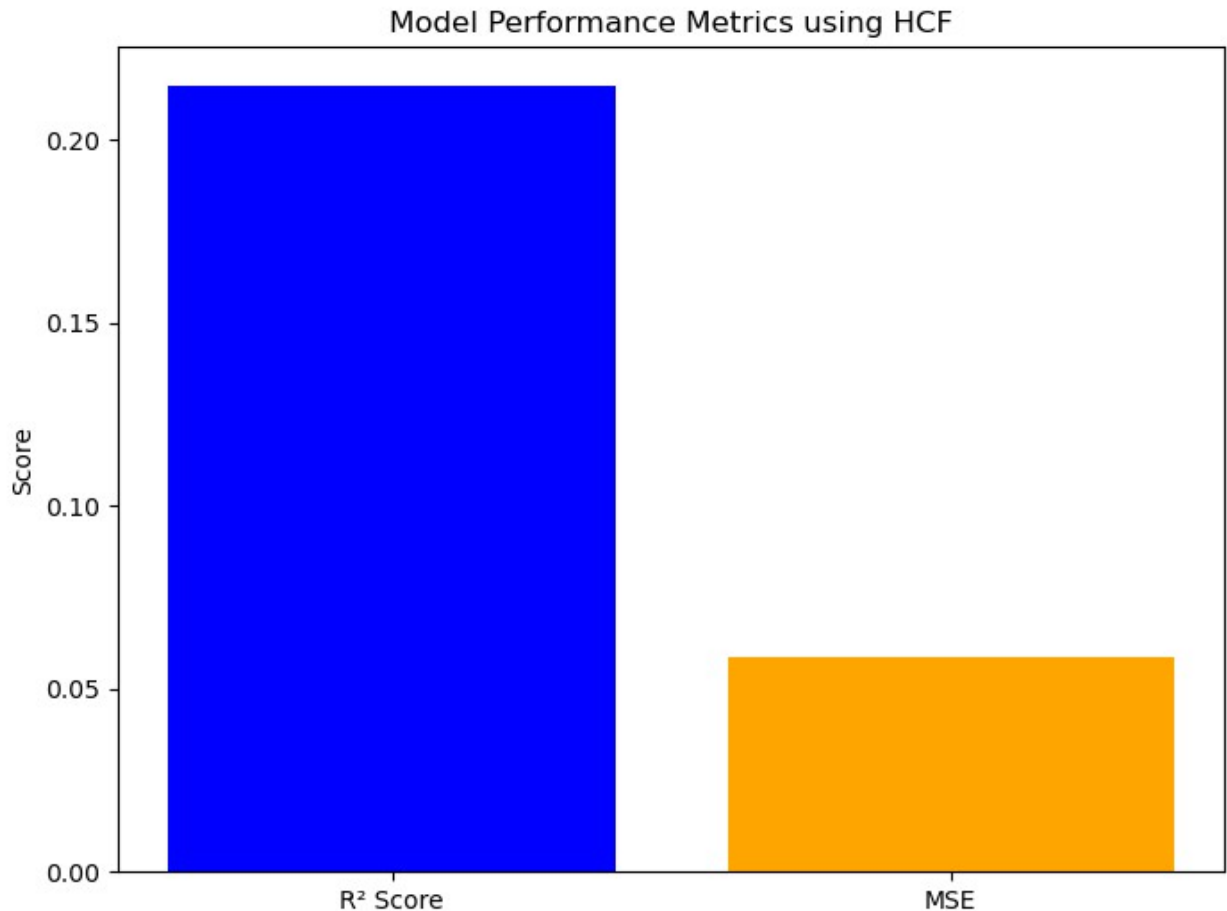
R2 Score: 0.21
Mean Squared Error: 0.06

# Example R2 Score visualization
metrics = {'R2 Score':r2LVF, 'MSE': mseLVF}
plt.figure(figsize=(8, 6))
plt.bar(metrics.keys(), metrics.values(), color=['skyblue', 'orange'])
plt.title("Model Performance Metrics using LVF")
plt.ylabel("Score")
plt.show()

#HCF# Example R2 Score visualization
metrics = {'R2 Score':r2HCF, 'MSE':msehcf}
plt.figure(figsize=(8, 6))
plt.bar(metrics.keys(), metrics.values(), color=['blue', 'orange'])
plt.title("Model Performance Metrics using HCF")
plt.ylabel("Score")
plt.show()

```





```
# Data for plotting
methods = ['Low Variance Filter (LVF)', 'High Correlation Filter (HCF)']
r2_scores = [r2LVF, r2HCF]
mse_values = [mseLVF, msehcf]

import matplotlib.pyplot as plt
import numpy as np

# Create a figure with subplots
fig, ax = plt.subplots(1, 2, figsize=(12, 5))

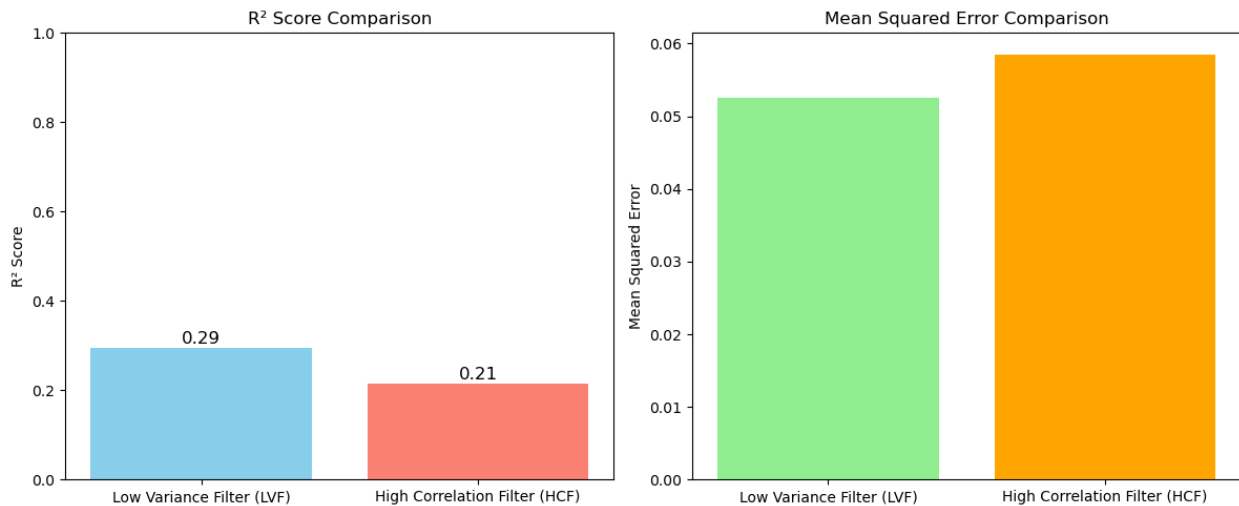
# Bar plot for R² Scores
ax[0].bar(methods, r2_scores, color=['skyblue', 'salmon'])
ax[0].set_ylim(0, 1) # R² Score ranges from 0 to 1
ax[0].set_title('R² Score Comparison')
ax[0].set_ylabel('R² Score')
for i, v in enumerate(r2_scores):
    ax[0].text(i, v + 0.01, f'{v:.2f}', ha='center', fontsize=12)

# Bar plot for Mean Squared Error (MSE)
```

```

ax[1].bar(methods, mse_values, color=['lightgreen', 'orange'])
ax[1].set_title('Mean Squared Error Comparison')
ax[1].set_ylabel('Mean Squared Error')
# Show the plots
plt.tight_layout()
plt.show()

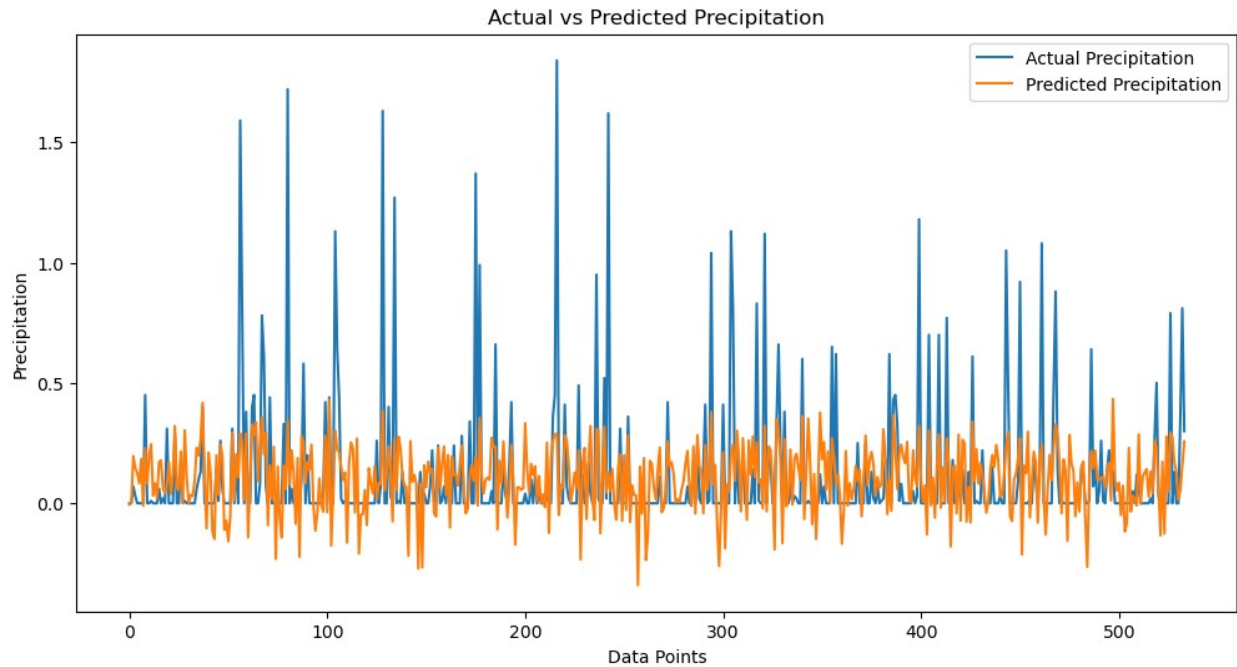
```



```

plt.figure(figsize=(12, 6))
plt.plot(y_test.values, label="Actual Precipitation")
plt.plot(y_pred, label="Predicted Precipitation")
plt.legend()
plt.title("Actual vs Predicted Precipitation")
plt.xlabel("Data Points")
plt.ylabel("Precipitation")
plt.show()

```



```
import seaborn as sns
plt.figure(figsize=(12, 8))
corr_matrix = X_selected.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title("Correlation Heatmap of Selected Features")
plt.show()
```

[illegible]